Blockchain and Trusted Computing: Problems, Pitfalls, and a Solution for Hyperledger Fabric

Marcus Brandenburger
IBM Research - Zurich

Christian Cachin
IBM Research - Zurich

Rüdiger Kapitza
TU Braunschweig

Alessandro Sorniotti
IBM Research - Zurich

ABSTRACT

A smart contract on a blockchain cannot keep a secret because its data is replicated on all nodes in a network. To remedy this problem, it has been suggested to combine blockchains with *trusted execution environments* (TEEs), such as Intel SGX, for executing applications that demand privacy. As a consequence untrusted blockchain nodes cannot get access to the data and computations inside the TEE.

This paper first explores some pitfalls that arise from the combination of TEEs with blockchains. Since smart contracts executed inside TEEs are, in principle, stateless they are susceptible to rollback attacks, which should be prevented to maintain privacy for the application. However, in blockchains with non-final consensus protocols, such as the proof-of-work in Ethereum and others, the contract execution must handle rollbacks *by design*. This implies that TEEs for securing blockchain execution cannot be directly used for such blockchains; this approach works only when the consensus decisions are *final*.

Second, this work introduces an architecture and a prototype for smart-contract execution within Intel SGX technology for Hyperledger Fabric, a prominent enterprise blockchain platform. Our system resolves difficulties posed by the execute-order-validate architecture of Fabric, prevents rollback attacks on TEE-based execution as far as possible, and minimizes the trusted computing base. For increasing security, our design encapsulates each application on the blockchain within its own enclave that shields it from the host system. An evaluation shows that the overhead moving execution into SGX is within 10%–20% for a sealed-bid auction application.

1 INTRODUCTION

Blockchain and distributed ledger technology (DLT) have received a lot of attention recently as means to distribute trust over many nodes in a network. The transparency and resilience gained from the decentralized protocol execution ensure the integrity of blockchain applications, or *smart contracts*, realized on such a "world computer." However, the proliferation of data on a blockchain directly contradicts the goal to keep the application state confidential and to maintain privacy for its users, a condition that exists for many intended applications of blockchain technology.

Although cryptographic protocols (such as secure multiparty computation and zero-knowledge proofs) offer attractive solutions for privacy on a blockchain, they are not yet mature enough to run general-purpose computations easily and to be widely deployed. As a promising alternative, the use of *trusted execution environments (TEEs)* for running blockchain applications has been proposed, especially by the industry working on *consortium blockchains*, where the consensus process is governed by controlled nodes [17, 32]. Intel's

Software Guard Extensions (SGX) is the most prominent TEE technology today and available together with commodity CPUs [2, 18, 31]. It establishes trusted execution contexts called *enclaves* on a CPU, which isolate data and programs from the host operating system in hardware and ensure that outputs are correct. An enclave might run only a small dedicated part of an application [11, 18, 29] or can contain an entire legacy system, including some operating-system support [5, 7, 44]. By using a TEE, one does not have to trust the host system of the enclave, which runs the blockchain node and participates in protocols.

For protecting smart contracts through TEEs, issues around rollback attacks, state continuity, and protocol integration for TEEs [42] must be addressed. As is well-known, there are non-trivial interactions between integrity violations and information leakage for stateful secure computation [1].

Motivating example. Imagine an auction of a digital item on a blockchain. In sealed-bid auction designs [34] (e.g., Vickrey auctions) keeping the bids secret is of primary importance, so that neither another bidder nor any other party can learn anything about them. Only a trusted auctioneer should learn the bids to the extent necessary for evaluating the auction. For moving the auction to a blockchain, the functions of the auctioneer are implemented by a smart contract. The distributed ledger stores encrypted bids such that the bidders are able to verify that their submitted bids were actually considered in the final evaluation. The blockchain nodes execute the auction's smart contract, which records the bids, closes the auction, evaluates it, and autonomously executes the transaction assigning the item to the winning bidder and transferring the payment to the seller.

By running the auction's code within an SGX enclave, the auction maintains privacy and simultaneously benefits from the transparency of the blockchain. More precisely, the bids are encrypted, the key to decrypt them resides only inside the enclave, and the smart contract controls operations with the key. The bidders commit their encrypted bids to the blockchain and the enclave decrypts them for determining the winner. However, this simplistic auction solution may leak information, as described next, whenever a malicious node can manipulate the operation invocation order.

State continuity and rollback attacks. As the industry is slowly realizing, rollback attacks on stateful applications running in TEEs pose serious risks, unless the *state continuity* of an application is ensured [9, 30, 37, 42, 43]. For instance, if a malicious blockchain node may influence the order in which transactions are executed by the enclave, the node can break the confidentiality of the sealed-bid auction even if it cannot decrypt the bids. In particular, the node

might cause the enclave to execute the evaluation transaction multiple times and reset the enclave again afterwards, every time when a new bid has been stored on the blockchain. Thereby the node could learn information about other bids. This illustrates how an integrity violation can lead to breaking confidentiality. (Although platforms like SGX provide access to non-volatile monotonic counters that might prevent rollback attacks, their use introduces considerable complications for tolerating crashes and they are often too slow [9, 37]. Hence, we do not consider them in this work.)

Rollback attacks can be prevented if the state input to the smart-contract enclave always corresponds to the unique, committed blockchain state. One way to guarantee the desired state continuity would be to run the whole blockchain node, especially its protocol logic and the state maintenance, within the enclave. This is often not feasible for practical reasons, however, and leads to other security issues because the code running inside the TEE has a large attack surface. However, in blockchain systems with *non-final* consensus protocols that may fork temporarily, a node remains prone to being rolled back *by design*, even when it resides completely in a TEE, because the underlying consensus protocol requires it.

Contributions. In this paper we examine the state-continuity problem for trusted execution on blockchains, arising from rollback attacks that malicious nodes might mount. We discuss why blockchains with consensus that has no final decisions, such as the "proof of work" in Bitcoin or Ethereum, are inherently unable to benefit from TEEs to maintain confidentiality. If the blockchain nodes hosting TEEs can access the final blockchain state in a trusted way, on the other hand, then such rollbacks can be prevented.

As the main contribution of this work, we design a secure solution for secure smart-contract execution on a blockchain using Intel SGX, the most prominent TEE technology available today, and *Hyperledger Fabric* [3], or *Fabric* for short, a flexible and modular platform for consortium blockchains. Fabric uses a modular notion of consensus whose outputs are always final, which avoids the protocol-inherent rollback attack. As Fabric is the most prominent technology for consortium blockchains today, our design can also be integrated with other, similar systems.

Some consortium blockchain platforms follow the generic approach to state-machine replication [38], where a consensus protocol first decides on an order among all transactions and the nodes *subsequently* execute them according to the decided order. In Fabric, however, the peers execute transactions and compute state updates *before* their relative order has been determined through a consensus protocol. The ordering process only uses the outcome of the transaction (i.e., the induced state changes) during consensus. While this offers a flexible programming model for smart contracts [3], it also introduces additional complications that must be considered.

We have implemented a prototype that enables smart-contract execution inside Intel SGX for Hyperledger Fabric. We demonstrate an auction application and evaluate the performance of the prototype compared to the unprotected execution. The results show that our prototype reaches 0.80x–0.90x the throughput of the unprotected implementation, which is acceptable for protecting the confidentiality.

Organization. This paper is structured as follows. In Section 2 we discuss why public blockchains with non-final consensus are inherently unable to execute smart contracts in TEEs. Section 3 introduces Intel SGX and Fabric, the two technologies used mostly in the remainder of the paper. In Section 4 we introduce the system model, describe the security goals, and discuss several approaches to run smart contracts in TEEs and their complications. Our solution to execute applications on Fabric with SGX is presented in Section 5 and its security is examined in Section 6. Performance evaluation results are reported in Section 7. Finally, Section 8 reviews related work and Section 9 concludes the paper.

2 CONSENSUS WITH NON-FINAL DECISIONS

Public blockchains patterned after Bitcoin [35] do not reach consensus with *finality*. Their consensus mechanism is based on a randomized protocol, in which for each epoch (or "block height") a node selected through a probabilistic scheme that is difficult to bias (such as a "proof of work") disseminates a block of transactions to be appended to the blockchain. Such blocks are propagated to all nodes with a peer-to-peer gossip protocol that is efficient but does not guarantee strict consistency. During regular operation, the view of the nodes in different parts of the network may diverge, and such forks are resolved through the protocol rule that the "longest" branch is adopted by all nodes as the valid blockchain and determines the state. As shown by multiple formal analyses of the protocol (e.g., [15]), the probability that such forks last for many epochs vanishes exponentially fast, but it cannot be made negligible for short forks. If the underlying network does not ensure universal connectivity, this can lead to devastating attacks on the safety of a public blockchain [4].

When a node first receives the block as a candidate that should extend the current chain, the node validates the block's content, including that all transactions inside are correct. For Bitcoin this validation is simply checking that a "coin state" has not been spent earlier, but for programmable blockchains like Ethereum [45], this entails executing all transactions and computing the corresponding state updates. If the block is valid, the node appends the block to its local chain and updates its state accordingly. But when the node later receives other blocks that are all valid and collectively extend a prefix of the currently held chain to a "longer" chain, the node reverts the earlier transactions and instead executes the subsequently received transactions. There is a significant probability that a node has to revert transactions during regular operation and, therefore, consensus is never final. In essence, a node must continue to participate in the consensus protocol forever, just to be sure that the blockchain state it holds remains valid. Blockchains using "proof-of-stake" consensus also suffer from similar forks (see the overview and analysis by David et al. [14]).

As becomes clear from this discussion, TEEs cannot be used to secure transaction execution and validation in blockchains based on non-final consensus. For example, an Ethereum virtual machine (EVM) running within SGX would have to produce the outputs resulting from a transaction immediately, but already during normal operation, the EVM could be rolled back to an earlier state that is beyond its control. This also holds if the consensus protocol is executed inside the TEE, as a malicious host controlling the communication could censor blocks from the network and forge valid blocks of its choice, given enough time. As argued before

through the auction example, application-level secrets could be revealed easily.

Therefore consensus with finality seems to be a necessary prerequisite to rely on TEEs for securing blockchains and for keeping transaction data secret. If one lets the TEE execute only transactions that are final, any attempt to roll back its state amounts to an attack, and such attacks can be prevented using existing methods for state continuity. This insight stands also behind some of the early designs and technologies that aim at this goal. For example, Microsoft's Coco Framework [32], available only as a white paper so far, uses the EVM but mentions quorum-based consensus with finality. In the Hyperledger Sawtooth platform (http://www.hyperledger.org/projects/sawtooth), which is most actively developed by Intel, the role of SGX technology lies in securing the "proof-of-elapsed-time (PoET)" consensus protocol, but SGX is not used for safeguarding secrets of a smart contract.

3 TECHNOLOGIES

In this section we review Hyperledger Fabric, an open-source blockchain platform developed under the Hyperledger Project (http://www.hyperledger.org) hosted by the Linux Foundation. We then describe Intel's Software Guard Extensions (SGX), which adds hardware-enforced security to the Intel CPU architecture and enables secure smart contract execution.

3.1 Hyperledger Fabric

Hyperledger Fabric [3] is a permissioned blockchain platform (run by a consortium), where multiple parties may participate and together form a distributed ledger network. The ledger records all interactions between the parties as transactions. A transaction invokes a smart contract called *chaincode*, which defines an application running on the blockchain.

A Fabric network consists of *clients*, *peers*, and an *ordering service*, as illustrated in Fig. 1. For each peer, a special client called *admin* has administrative control over the peer, for instance in order to install a chaincode. The basic transaction flow is as follows: (1) A client invokes a chaincode by sending a transaction proposal to one or more peers, which (2) execute the chaincode and produce a proposal response called *endorsement*. (3) The client then collects the endorsements and assembles them to a transaction that it submits to the ordering service. (4) The ordering service establishes the total order of all transactions and broadcasts them as blocks of transactions to all peers in the network. When a peer receives a block, it validates every transaction, eliminates those that were based on state that has become invalid, and commits the valid ones to its local ledger.

Other blockchain platforms execute transactions after ordering them [13], e.g., JPMC Quorum (https://github.com/jpmorganchase/quorum), Hyperledger Sawtooth (https://github.com/hyperledger/sawtooth-core), or Chain Core (https://chain.com/). In contrast, Fabric uses a three-phase *execute-order-validate* architecture. In the remainder of this section we provide more details of each phase.

Chaincode execution and endorsement. A transaction of a chaincode is executed by a set of *endorsing peers* for the chaincode during the endorsement phase. Initially the chaincode is installed on every endorsing peer by an admin. The clients invoke the chaincode by


Figure 1: A high-level view of a Fabric network, illustrating the transaction flow invoked by a client (left).

sending a transaction proposal containing a chaincode operation to the peer. The chaincode takes the operation as input, processes it according to its smart-contract application, and may return an optional execution result.

While executing, the chaincode may access the blockchain state, which is provided as a *key-value store (KVS)*, with *getState* and *putState* operations. The *putState* operation does not immediately update the peers' local state; instead, it records the change in a *writeset*, containing the updated keys and their values. Additionally, all keys accessed by the transaction during execution and their versions (i.e., the positions in the history where they were last updated) are collected in a so-called *readset*.

Taken together, the execution result, the readset, and the writeset form a transaction proposal response or *endorsement*. The peer returns this to the client. Through executing the chaincode and producing the endorsement, the peer vouches for the correct execution of the chaincode and *endorses* the change. An endorsement policy specifies the endorsing rules for each chaincode (e.g., who are the endorsing peers or how many endorsements are needed). Accordingly, the client collects sufficiently many endorsements and integrates them to a transaction that it submits for ordering.

Ordering. The ordering service in Fabric is responsible for establishing the *total order* of all transactions in the blockchain and therefore to ensure a consistent view of all transactions across all peers. Typically, the ordering service consists of multiple nodes for scalability and resilience, and leverages a protocol to reach consensus on the total order. Clients submit transactions created in the endorsement phase to the ordering service. For efficiency, transactions are distributed among all peers in batches or *blocks*, using a gossip protocol [6].

Validation and state updates. A peer receives a block of transactions from the ordering service and utilizes a *validation system chaincode (VSCC)* to validate each transaction and apply its effects to the local ledger. Validation is a deterministic process and performed by every peer. In particular, the peer checks that every transaction fulfills the endorsement policy. Then, a *read-write conflict* check is performed, that is, the peer verifies that the versions in the readset match the current blockchain state at the peer. If both validation checks are successful, the updates in the writeset of the transaction are applied to local blockchain state of the peer. An invalid transaction has no effect on the state and the issuing client should reinvoke the transaction again. A Fabric blockchain is initialized through a so-called *genesis block*, which is created collaboratively by all network participants.

3.2 Trusted execution with Intel SGX

Modern trusted execution environments (TEEs), such as Intel's Software Guard Extensions (SGX) [2, 18, 31], add hardware-enforced security to commodity platforms. SGX enables Fabric peers to execute chaincode in a trusted execution context, also called *enclave*. Particularly, an enclave defines an isolated memory area that is guarded by hardware-enforced mechanisms, which guarantee *confidentiality* and *integrity* of an enclave even if the entire platform is compromised. That is, even higher-privileged code (e.g., the operating system) can neither access that memory area nor modify it without being detected.

Enclave protection and attestation. SGX enforces that only genuine applications are executed in an enclave. "Genuine" means that the code has not been tampered with and operates precisely as intended by the developer. For this reason, a cryptographic hash called *mrenclave* of the code and data initially loaded into the enclave is generated by the CPU. If *mrenclave* matches the hash of the genuine application signed by the developer, the enclave starts successfully. This ensures that the correct application (e.g., a specific chaincode) is runs in an enclave.

During runtime, an enclave is capable to prove to a third party (e.g., a Fabric client or another peer) that a specific application is loaded and executed in an actual enclave on a SGX platform. For this purpose a procedure called *remote attestation* [2] is used. It works as follows: A client with prior information about *mrenclave* of the target enclave sends an attestation challenge to the enclave host (e.g., the peer) and in return receives a proof ϕ, also known as *quote*, produced by a target enclave and the platform. We use the term *attestation report* to refer to a quote. The client forwards ϕ to the Intel Attestation Service (IAS), which verifies it using a *group signature* scheme called EPID [12], and replies with an attestation result. The attestation result either confirms that ϕ was actually produced by an SGX enclave running the intended code or indicates that ϕ is invalid. The enclave can also embed custom data in ϕ, which builds the basis for key exchange protocols.

Remote attestation involves an intermediate step called *local attestation* that is performed between two enclaves on the same platform. In this step, a special enclave called *Quoting Enclave (QE)* verifies that ϕ was produced by an enclave on the same platform, using HMAC with a shared secret key only accessible by enclaves on the same platform. Local attestation can be performed by any two enclaves on the same platform.

Enclave state and data sealing. Since enclaves reside in a protected memory area in the CPU, enclaves are volatile, thus, when an enclave stops, restarts, or just crashes, its internal state is lost and cannot be recovered. For this reason SGX supports data sealing, a mechanism that allows to encrypt and authenticate data before it leaves an enclave and is stored externally (e.g., on persistent storage). After an enclave restarts, it may load the sealed data and decrypt it. Although the sealing mechanism protects data confidentiality, it does not prevent rollback attacks, that is, an attacker may cause an enclave to recover from properly sealed but stale data [9, 30, 43]. For many applications this poses serious problems and they must be protect against it, however, for stateless applications, such as chaincodes, this is not relevant.

Moreover, enclaves have access to a secure random number generator that allows to build cryptographic primitives, such as key generation, encryption, and digital signatures. Both, remote attestation and data sealing rely on a cryptographic key-management infrastructure rooted in a secret key fused into the CPU, which provides deterministic key-derivation functions to an enclave.

4 PROBLEM DESCRIPTION

This section describes the problem of secure smart-contract execution using trusted hardware for blockchains with *final* consensus, in particular, Hyperledger Fabric [3]. We explore intricacies that may still be caused by rollback attacks in this setting, illustrate a strawman approach that is infeasible, and introduce our approach to support secure chaincode execution using Intel SGX. This executes each chaincode in its own enclave during endorsement at a peer and thereby protects the confidentiality and integrity of the blockchain application.

4.1 System model

We consider a Fabric blockchain network with clients, an ordering service, and a set of peers, which collaboratively execute transactions and maintain a distributed ledger on a single Fabric "channel."

A client invokes transactions by sending a chaincode operation to some peer, which then executes (simulates) it and produces an endorsement containing the resulting state change on the ledger. The operation, the response, as well as the ledger may contain sensitive information that should stay secret.

To prevent such information leakage, every peer is equipped with an SGX-enabled CPU and executes transactions inside an enclave. The chaincode is stateless, and a transaction only takes the operation and the blockchain state in the KVS as inputs, accessed with *getState*. The chaincode must perform updates to the ledger only through *putState* operations. The execution of a chaincode operation returns a response that may include a computation result, the state update, and the read-write dependencies.

4.2 Threats

Although most peers are usually correct, a peer may become *malicious* and behave incorrectly, for instance, when it tries to maximize its own profit or becomes corrupted by an attacker. A peer has full control over the operating system, applications, and the data residing in memory and persistent storage (i.e., the blockchain state). A malicious peer, however, cannot access or tamper with the code and data residing in an enclave (see Section 3.2). A malicious peer may neither break cryptographic primitives nor extract any secret information that from an enclave. Consequently, a chaincode running in an enclave always produces the correct results, that is, the chaincode does not deviate from its specification, the enclave-internal state is only known to the enclave itself, and nothing is revealed apart from the resulting state change.

However, a malicious peer can invoke the chaincode enclave with any input and in arbitrary order. The peer may intercept, modify, reorder, discard, or replay chaincode operations, and when the chaincode enclave accesses the KVS, the peer may feed any blockchain state to it.

Furthermore, the peer might even drop messages or completely halt an enclave, but we do not consider such denial-of-service attacks in this work. We also ignore potential information leakage from SGX on side channels [10, 39, 41, 46] because this appears orthogonal to our focus.

As is well-known from the literature on secure computation with cryptographic protocols [1], integrity and confidentiality cannot be considered separately. Likewise, for a secure application running in an enclave, a malicious host may break confidentiality by triggering the enclave to execute on "incorrect" inputs. In the blockchain context, this means that the chaincode execution deviates from the consensus-based transaction order.

Repeating and extending the auction example from the introduction, such an attack could reveal secret information as follows. Suppose that evaluating the auction on the current blockchain state s_1 would let a bid b_1 win the auction. If the malicious node can trigger the auction-evaluation transaction, it learns b_1. If the node can reset the enclave to s_1 and execute another transaction, it can submit a bid b_2, add it to the ledger, subsequently evaluate the auction, and learn if $b_2 > b_1$. Such a rollback attack clearly breaks the confidentiality of the individual bids. As mentioned earlier, rollback attacks on trusted execution environments and their prevention has only recently been understood better [9, 30, 43].

4.3 Strawman approach

It follows from the discussion in Section 2 that letting the enclave only execute transactions that have been ordered by the network with finality prevents the rollback problem. This amounts to running the entire blockchain peer inside an enclave, as also suggested by Microsoft Coco [32, 40] and related work. We call this the *strawman approach* that might work for an order-execute architecture where the consensus process has only final decisions, but we argue later why better designs exist.

For Fabric, the strawman design would mean to encapsulate the chaincode execution, endorser, committer, ledger-state access, and all other parts of a peer inside an enclave. This obviously protects the integrity of the input sequence for the chaincode, since the entire Fabric peer runs within SGX. A similar approach is taken in the blockchain-as-a-service platform of IBM, which deploys Fabric peer as a *secure service container* on an IBM Z system. The secure container includes the whole operating system, middleware stack, and blockchain platform [19].

Although no operating system is running within SGX, recent research has demonstrated how legacy applications can run in SGX through a library OS that executes unmodified applications in an enclave [5, 7, 44]. Note that the library OS adds tens of thousands of lines of code that also run along the application in the enclave. This approach introduces multiple problems, however. First, it stands in contrast to the important computer-security principle of minimizing the size of the *trusted computing base (TCB)*. Specifically, also the SGX developer guidelines [21] recommend to partition an application into a trusted and an untrusted component; only a small portion of the application code should execute inside the enclave. A smaller TCB has fewer errors, reduces the attack surface, and is more amenable to security analysis than the entire application.

A second problem stems from the limited memory available to enclaves. An enclave's memory resides in the *enclave page cache (EPC)* isolated from the rest of the system. The EPC is currently limited to 128 MB. Once an enclave reaches that limit pages are outsourced to DRAM. This results in a dramatic loss of performance, as reported in several works [5, 11, 36]. In particular, since the ledger grows with every block, holding the whole blockchain state in the enclave quickly reaches the memory limitation.

4.4 Approach for Hyperledger Fabric

To avoid the drawbacks of the strawman approach, we adopt a modular architecture that separates the chaincode execution conceptually from the peer and moves the execution into an enclave. The protocol-specific aspects of the peer are encapsulated in an abstract ordering service, of which one process might run on the same peer. The ordering service is *trusted* in the sense that it cannot be rolled back.

The ordering service produces a signed sequence of transactions for execution within the enclave. The enclave can verify that transactions originate from the ordering service, are in the proper order, and have not been tampered with. The enclave also keeps information about the transaction history, which allows to detect transaction-ordering violations or replayed transactions. The malicious host might still reset the enclave to an earlier point in the execution sequence, but this would not harm the application since the transactions are deterministic and execution would simply produce the same outputs again.

As described so far, this approach works fine with an order-execute architecture for state-machine replication. Fabric, however, uses the execute-order-validate paradigm, where a peer executes a transaction before consensus on the order is reached (see Sec. 3.1). Consequently the execution is speculative and can be repeated without affecting the blockchain state, as transactions are simulated during endorsement and only take effect after the ordering. This means a malicious host could infer information about the secret application data from the speculative execution. Not even a trusted ordering service can prevent this type of leakage.

To resolve this issue, we will have to adapt the applications to respect the speculative nature of execution in Fabric. For the auction example, in particular, a *barrier* will be stored on the blockchain such that the chaincode enclave only evaluates the auction if the barrier is present. The barrier is set by invoking the chaincode with a transaction to "close" the auction but not yet evaluate it. If the barrier is present on the ledger, a malicious peer may no longer submit new bids to the auction. On the other hand, the auction evaluation will only consider bids added to the ledger before the barrier. Note that this barrier plays a role similar to a memory barrier in a multi-core computer system with concurrent threads.

Following the execute-order-validate architecture, the chaincode enclave must execute transactions only on the committed blockchain state, that is, with ledger entries that result from ordered transactions and that have been committed by all peers. Otherwise a malicious peer may produce the barrier itself and feed the resulting state into the enclave when evaluating the auction. The system described in the next section ensures this.

To formally model the information leakage permitted in the execute-order-validate architecture of Fabric, we model a blockchain

as stateful functionality $F : \mathcal{S} \times \mathcal{T} \rightarrow \mathcal{S}$. At any time the state of the chaincode is an element of $\mathcal{S}$. The clients invoke transactions in $\mathcal{T}$, which may contain operations with arguments according to F, but these are subsumed into the different $t \in \mathcal{T}$. Given $s \in \mathcal{S}$, applying a transaction $t \in \mathcal{T}$ of F means to compute $s' \leftarrow F(s, t)$, resulting in a subsequent state $s' \in \mathcal{S}$. Using a trusted ordering service as introduced earlier, the blockchain's state evolution is defined through the sequence of transactions signed by ordering.

With the chaincode functionality F running in an SGX enclave, even a malicious peer may only learn the subsequent state resulting from a transaction, but nothing about the computation itself. Since cryptographic keys could reside in the enclave, the ledger state doesn't necessarily reveal all relevant information. Due to the rollback attacks introduced earlier, however, such a peer can execute any transaction *on any input state* that is in the history of transactions issued by the ordering service.

Definition 4.1 (Security up to resets). Consider a blockchain system with an execute-order-validate architecture and suppose the correct ordering service produces a sequence of states $\langle s_0, s_1, \ldots, s_m \rangle$, where $s_j = F(s_{j-1}, t_j)$ for $t_j \in \mathcal{T}$ and $j \in [1, m]$. We say that the chaincode is *secure up to resets* if any malicious peer, through interacting with the chaincode running inside the enclave, may obtain states $s^*_{k+1} = F(s_k, t^*)$, for any $k \in \{0, 1, \ldots, m\}$ and an arbitrary transaction $t^* \in \mathcal{T}$, but no further information.

The security-up-to-resets notion formalizes attacks on TEE-based execution in Fabric, where a malicious peer may collude with a client. The client invokes an arbitrary transaction t^* that reveals information about the chaincode's state. The peer lets the TEE execute t^* and produce an output, but the endorsement is never sent for ordering and the transaction is never appended to the blockchain. The chaincode may leak all states resulting from such executions.

Note that Fabric permits parallelism during execution for separating trust assumptions and for increasing scalability. By adding a barrier into the blockchain, an application essentially benefits from the guarantees of the order-execute design with respect to rollbacks across the barrier. Requiring a barrier after every transaction would actually impose the order-execute paradigm onto Fabric.

5 SECURE CHAINCODE EXECUTION

This section describes our system for secure chaincode execution in Hyperledger Fabric using Intel SGX.

5.1 System architecture

Our approach extends a Fabric peer with the following components: A *chaincode enclave* that executes a particular chaincode and a *ledger enclave* that enables all chaincode enclaves to verify the blockchain state integrity; all run inside SGX. In the untrusted part of the peer, an *enclave registry* maintains the identities of all chaincode enclaves and an *enclave transaction validator* that is responsible for validating transactions executed by a chaincode enclave before committing them to the ledger. Fig. 2 shows the components.

Chaincode enclave. The chaincode enclave executes one particular *chaincode*, and thereby isolates it from the peer and from other chaincodes. A *chaincode library* acts as intermediary between the
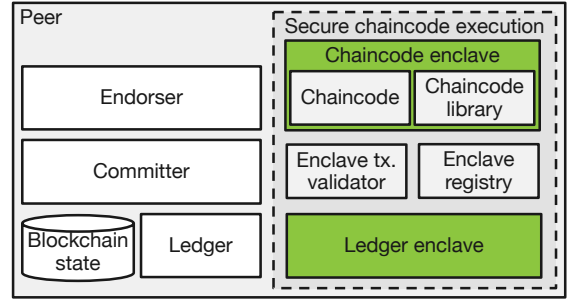

Figure 2: System architecture. The dashed box denotes the components added to the peer to enable secure chaincode execution with SGX. The components running within SGX enclaves are denoted in green (or dark) color.

chaincode in the enclave and the peer. The chaincode library exposes the Fabric chaincode interface and extends it with additional support for state encryption, attestation, and secure blockchain state access.

Ledger enclave. The ledger enclave maintains the ledger in an enclave in the form of integrity-specific metadata representing the most recent blockchain state. It performs the same validation steps as the peer (see Sec. 3.1) when a new block arrives, but additionally generates a cryptographic hash of each key-value pair of the blockchain state and stores it within the enclave. The ledger enclave exposes an interface to the chaincode enclave for accessing the integrity-specific metadata. This is used to verify the correctness of the data retrieved from the blockchain state.

Enclave registry. The enclave registry is a chaincode that runs outside SGX and maintains a list of all existing chaincode enclaves in the network. It performs attestation (see Sec. 3.2) with the chaincode enclave and stores the attestation result on the blockchain. The attestation demonstrates that a specific chaincode executes in an actual enclave. This enables the peers and the clients to inspect the attestation of a chaincode enclave before invoking chaincode operations or committing state changes.

Enclave transaction validator. The enclave transaction validator complements the peer's validation system and is responsible for validating transactions produced by a chaincode enclave. In particular, the enclave transaction validator checks that a transaction contains a valid signature issued by a registered chaincode enclave. If the validation is successful, it marks the transactions as valid and hands it over to the ledger enclave, which crosschecks the decision before it finally commits the transaction to the ledger.

5.2 System initialization

When a peer joins the blockchain network, the ledger enclave is initialized by the admin with the genesis block, which contains the blockchain configuration and the expected hash (*mrenclave*) of the ledger enclave. If the actual *mrenclave* obtained by the peer does not match the value in the genesis block, the ledger enclave does not proceed with the initialization. The ledger enclave then generates a private/public key pair (SK_{LE}, PK_{LE}), which allows to uniquely identify the ledger enclave. The public key is revealed

to the chaincode enclaves whereas the private key is kept secret within the ledger enclave.

The ledger enclave maintains several configuration values initially obtained from the genesis block, such as the identities (i.e., public keys) of the peers, the clients, and the ordering service, which are used to authenticate all received blocks and transactions. The ledger enclave only accepts blocks that come from the ordering service as defined in the genesis block. To ensure this, it verifies that each block has a valid signature issued by the ordering service. Note that the blockchain consortium configuration can be updated using configuration transactions. For simplicity, however, we assume a static consortium.

Every block has a sequence number and contains a list of transactions. The ledger enclave maintains information about the most recently processed transaction, to ensure that all blocks are processed in the correct order and no blocks are missing.

Once the peer has joined the network and has started its ledger enclave, the peer admin also installs the enclave registry on every peer and instantiates it. This completes the initialization of the peer.

5.3 Chaincode enclave bootstrapping

We now describe how to initialize a chaincode enclave. This is initiated by the peer admin and consists of the following phases: (1) creating the chaincode enclave; (2) registering with the enclave registry; (3) provisioning of secrets; and (4) binding the chaincode enclave to the ledger enclave. These phases are explained next.
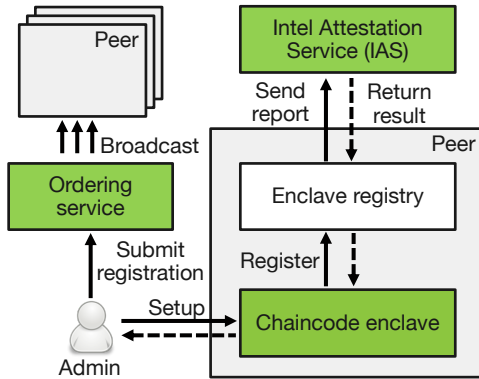

Figure 3: Enclave registration process.

In the first phase, the admin installs the chaincode enclave at the peer and then sends a SETUP transaction proposal. The peer then starts the chaincode enclave, which generates a private/public key pair (SK_{CC}, PK_{CC}). As for the ledger enclave, the public key is used to uniquely identify the chaincode enclave.

Second, the chaincode enclave registers itself with the enclave registry as shown in Fig. 3. For this purpose, the chaincode enclave calls REGISTER and in turn, the enclave registry performs remote attestation of the chaincode, as described in Section 3.2. In detail, the chaincode enclave first produces an *attestation report* that manifests that it is properly instantiated with a specific chaincode and it is identified by PK_{CC}. The report contains $mrenclave_{CC}$ and (hash of) PK_{CC}. The chaincode enclave then calls REGISTER at the enclave registry with the report and its public key as arguments. The enclave registry first checks that the report contains the expected

$mrenclave_{CC}$ and the correct hash of PK_{CC}. Subsequently, it sends the report to the IAS for verification and in return receives an *attestation result*, which shows whether the report was valid or not. Note that the attestation result is signed by the IAS and its verification key is publicly available. If the verification succeeds, the enclave registry completes the registration by calling *putState* to store the attestation result together with PK_{CC} on the ledger. This makes the attestation result accessible to all peers in the network through the ledger, certifying that this enclave runs the particular chaincode on the given peer. Clients and other peers use this in two ways. First, a client verifies that it invokes transactions involving secret data on an enclave authorized for this. Second, the enclave transaction validator of a peer, which updates the blockchain state, verifies that the execution results are genuine and result from the secure execution in the enclave.

After successfully registering the chaincode enclave, the admin optionally provisions the chaincode enclave with secrets. For instance, the admin may inject an encryption key for data stored on the blockchain into the chaincode enclave (see also Sec. 5.5).

In the last phase, the chaincode enclave binds to the ledger enclave through local attestation (see Sec. 3.2). This means that the chaincode enclave requests the ledger enclave to prove that it is runs the expected ledger-enclave code and runs on the same host platform. The ledger enclave produces an attestation report and returns it to the chaincode enclave, which then performs the same verification steps as described above. (In contrast to remote attestation, the cryptographic protection of local attestation uses HMAC and a shared key for verification, provided by the SGX platform.) If the verification succeeds, the chaincode enclave stores the ledger enclave's public key PK_{LE} and thereby binds itself to the ledger enclave, in the sense that the chaincode enclave uses this for verifying accesses to the blockchain state. The chaincode enclave rejects any blockchain state values not originating from this ledger enclave.

5.4 Chaincode execution

Endorsement. A client triggers the chaincode execution by sending an INVOKE transaction proposal with a chaincode operation to the peer. The peer forwards the chaincode operation to the chaincode enclave, which then processes it according to the smart contract. The chaincode enclave prepares a response and returns it to the peer, which subsequently sends it as a a transaction proposal response to the client.

In more detail, prior to invoking the chaincode enclave, the client queries the peer to retrieve the enclave's public key PK_{CC} and the corresponding attestation result from the enclave registry. The client then verifies the authenticity of the attestation result, using the IAS verification key, and checks that the attestation contains the expected *mrenclave* of the chaincode enclave, matching PK_{CC}. If the verification succeeds, the client invokes the chaincode enclave by preparing a transaction proposal for the target chaincode. In particular, the client encrypts the chaincode operation using PK_{CC}, and then sends the proposal to the peer, which extracts the chaincode operation and relays it to the chaincode enclave. Inside the enclave, the chaincode library decrypts the operation using SK_{CC} and invokes the chaincode with the operation as argument.

The chaincode processes the operation, produces a result, and returns it to the chaincode library. The chaincode may access the blockchain state using the chaincode library, which performs verifies the accesses as described in Section 5.5.

To complete the chaincode invocation, the enclave library creates a response, signs it using SK_{CC}, and returns it to the peer. The response includes the operation, the readset and the writeset, and the execution result. Optionally, the chaincode library encrypts the execution result before it leaves the enclave using an encryption key provided by the client. The peer then sends the transaction proposal response back to the client, which outputs the execution result, and submits the transaction to the ordering service.

Validation and state update. The ordering service accepts transactions submitted by the clients, assigns them to a block, and broadcasts the block to all peers in the network. In order to finalize a transaction, every peer validates the transaction and updates its ledger copy.

For validating transactions produced by a chaincode enclave, the enclave transaction validator essentially performs the same steps as the validation system chaincode (VSCC), checking for conflicts and evaluating the endorsement policy (Sec. 3.1). Additionally it verifies that the transaction was produced by the correct chaincode enclave as follows. The validator accesses the enclave registry to retrieve the attestation result and public key for the enclave indicated by the transaction. Then it verifies these following the same steps as described earlier. Subsequently, it also verifies the enclave's signature on the transaction. If this succeeds, the enclave transaction validator marks the transaction as valid, the peer commits the transaction to its local ledger, and updates the blockchain state accordingly.

5.5 Accessing the blockchain state

Recall that a chaincode in Fabric must only use and access state on the blockchain. The chaincode library together with the ledger enclave protects this data from manipulation by the local peer.

State integrity and consistency. As illustrated in Fig. 4, when the chaincode calls *getState(k)* to access the data for key k, the chaincode library loads the corresponding value *val* from the blockchain state in the enclave through chaincode API provided by the peer. Additionally, the chaincode library requests the corresponding integrity metadata from the ledger enclave by calling *getMeta(k, z)* with a nonce z. The ledger enclave returns the expected hash h_{val} of *val* in the blockchain state and a signature ϕ, produced by the ledger enclave as $\phi = sign_{SK_{LE}}(k\|z\|h_{val})$. The chaincode library has obtained PK_{LE} during bootstrapping and uses this to verify ϕ. If the signature verification over $k\|z\|Hash(val)$ succeeds, then *val* is correct according to the state of the ledger enclave. The nonce ensures that the response is fresh.

State confidentiality. The chaincode library may also protect the confidentiality of the blockchain state maintained by the chaincode enclave. The native data sealing method of SGX for protecting persistent data (Sec. 3.2) is not suitable for data shared by multiple enclaves on different peers. The reason is that sealed data can only be unsealed again by the same enclave.
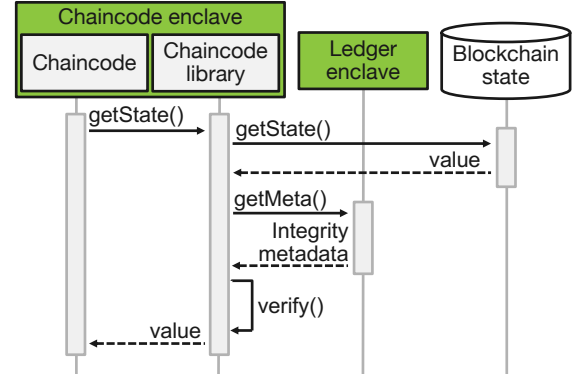

Figure 4: Blockchain state verification with the help of the ledger enclave.

Instead, the chaincode library provides a state-encryption mechanism that supports two modes: *client-based encryption* and *encryption per chaincode*. With client-based encryption, the client is responsible for key management and must provide an encryption key together with each chaincode operation. For chaincode-based encryption, a chaincode-specific key must be provisioned by an admin to all chaincode enclaves during bootstrapping. In both modes, encryption and decryption occur transparently to the chaincode during the *putState* and *getState* calls, respectively.

As an additional benefit of client- or chaincode-based key management compared to SGX-native methods, data on the blockchain can also be retrieved from the blockchain later, without the support of an enclave.

5.6 Reboot and recovery support

A system crash or reboot terminates all enclaves instantiated on the peer. In order to tolerate these without manual intervention, the internal states of each chaincode enclave and the ledger enclave are stored on persistent storage periodically.

The ledger enclave leverages sealing (Sec. 3.2) to protect its state (including integrity metadata and private key). For ensuring state continuity across crashes, the peer, in principle, has to write the ledger-enclave state to disk synchronously after each block has been processed. This clearly impacts the performance and can be mitigated in practice by defining a block interval for persisting the enclave state.

The chaincode enclave, in contrast, only has immutable state (including its private key) that was created during the enclave bootstrapping. It is sufficient to seal and store this once, after initialization. When the chaincode enclave restarts and restores itself from the sealed state, it will retain the same enclave identity and does not need to perform registration or remote attestation again.

5.7 Extensions

Support for confidential chaincode. Our chaincode enclave can be also extended to support the execution of confidential chaincode, which requires support for dynamic loading of encrypted code in enclaves [16, 44]. This allows to deploy proprietary smart-contract code without revealing it to the executing peers.

To enable this feature, the chaincode enclave is extended with a bootloader that injects an encrypted chaincode binary in the enclave. The bootloader then decrypts it and executes the chaincode. For this purpose, the admin that installs the chaincode on a peer, or the chaincode developer, encrypts the chaincode binary for a the chaincode enclave using its public key.

Furthermore, the attestation functionality of the chaincode enclave has to be adapted so that the peers and clients can verify that a specific, encrypted chaincode is executed by the enclave. Since *mrenclave* denotes the bootloader code running in the enclave, the attestation must also contain a hash of the chaincode binary, which is publicly known by the peers and clients.

Trusted state transfer. When a peer joins an existing blockchain it has to validate all blocks processed before and reconstruct the current blockchain state. Depending on the age of the blockchain, this effort may be prohibitive. Also, a peer might have been offline for a longer time and must catch up when it comes online again.

If the peer does not want to trust another peer for providing it the most recent blockchain state, then it can leverage a ledger enclave to obtain the current state securely.

When peer P_A joins or resumes after a long intermission, it contacts another peer P_B for support. P_A sends a message containing the hash of the genesis block and integrity metadata dependent on the position of its ledger enclave LE_A on the blockchain. Peer P_B passes this to its ledger enclave LE_B, which performs four steps: (1) It checks that LE_A is part of the same blockchain by comparing the hashes of the genesis blocks; (2) it calculates the difference Δ (in terms of KVS keys) between P_A's state and its own state from P_A's integrity metadata; (3) it creates an attestation report containing Δ and its last known block sequence number; and (4) it returns the report and Δ to P_B. At this point Δ only contains KVS keys and the corresponding integrity metadata, thus, P_B complements Δ with the actual values from the blockchain state. P_B also sends the attestation report to the IAS for verification. Next, P_B sends the attestation result and Δ to P_A, and P_A verifies its contents. If successful, then P_A updates its last known block sequence number and its blockchain state accordingly, and passes the data to LE_A, which performs the same verification steps as the peer and also updates the integrity metadata.

6 SECURITY

This section argues that the secure chaincode execution system presented in the previous section preserves security up to resets. Recall that this security notion is defined with reference to a sequence of blockchain states produced by transactions as decided by the trusted ordering service O. For being secure up to resets, any set of malicious peers interacting with SGX TEEs that host a chaincode CC deployed on Fabric must not be able to infer more than what is given from any transaction of CC invoked on one of these states. Our informal argument proceeds in three steps.

1. Any state update (in the form of a writeset) produced by a chaincode enclave with a public key PK_{CC^*} and accepted by a ledger enclave into the state of its peer originates from an enclave whose attestation report is stored on the ledger with public key PK_{CC^*}. Furthermore, any transaction output produced by a chaincode enclave,

for which a correct client has successfully attested the output to an enclave with key PK_{CC^*}.

This follows from the operations of the enclave registry and the enclave transaction validator. In particular, the enclave registry performs remote attestation with the chaincode enclave and thereby creates the attestation report that it stores in the ledger. This convinces the clients and the peers that the chaincode enclave has been instantiated with the chaincode represented by the *mrenclave$_{CC}$* value in the ledger. The correct clients and peers obtain their state in the form of the sequence of blocks with updates from O and can verify the integrity of the state updates signed by PK_{CC^*}.

2. On any peer, the ledger state entries obtained by chaincode CC inside an enclave represent the blockchain state after executing a prefix of the sequence of valid state updates that are output by O.

Note that a malicious peer may reset the ledger enclave at will to one of the sealed and persistently stored states that the enclave produces. Due to the VSCC checks and the monotonically increasing sequence of block numbers that the ledger enclave expects from O, the blockchain state represented within the ledger enclave always results from executing the sequence of transactions determined by O and deemed valid by VSCC and the endorsement policies. When the chaincode inside the chaincode enclave accesses state in the KVS, the mutual authentication between the ledger and the chaincode enclave, and the blockchain state-verification mechanism in Section 5.5 ensure that the state entries obtained by CC are correct according to the state of the ledger enclave. Since the ledger enclave holds the state after executing a prefix of the transaction sequence from O, the above statement follows.

3. Any state held by a chaincode within an enclave remains confidential up to what is revealed by executing transactions of the chaincode, invoked on a prefix of the complete sequence of valid state updates that are output by O.

This holds because the enclaves' execution logic and data are protected within the TEE. Contents of the ledger enclave are sealed before they are written to persistent storage, hence they cannot be altered by a malicious peer without being detected. The state of the chaincode enclave itself remains unchanged after initialization and is stored by the peer. However, all correct peers verify that they only interact with chaincode enclaves registered on the ledger itself. The ledger enclave may also contain an encryption key for protecting data on the ledger through the chaincode library, which handles state encryption and decryption transparently.

7 EVALUATION

We have built a prototype of the design for secure chaincode executing using Intel SGX with Hyperledger Fabric. This section describes the implementation and reports on the evaluation of the SGX prototype for the blockchain auction application.

7.1 Implementation

We implemented the prototype on top of Hyperledger Fabric 1.0. Each component of the architecture in Section 5.1 has been integrated with the Fabric peer code. We use the Intel SGX SDK for Linux 2.1.2 [22] to implement the components residing in an enclave such as the chaincode library and the chaincode. These components are written in C/C++. The other components, such as the untrusted

part of the chaincode enclave, the enclave registry, and the enclave transaction validator are written in Go. The ledger enclave runs as a system chaincode, which allows it to be integrated in the peer without major changes. This means that the chaincode enclave can access the ledger enclave via *chaincode-to-chaincode (cc2cc)* invocations. The ledger enclave uses a simple KVS based on *std::map* to store the integrity metadata.

The prototype supports a subset of the original fabric chaincode shim and provides *getState*, *getRangeQuery*, and *putState* calls. Those functions are implemented with the help of the untrusted part of the chaincode enclave and manage the data in the KVS on persistent storage. When the chaincode accesses the KVS to retrieve a value, the prototype contacts the trusted ledger enclave via a cc2cc invocation in order to retrieve the corresponding integrity metadata. This integrity protection uses HMAC-SHA256 with a verification key generated during bootstrapping (Sec. 5.3).

The enclave registry interacts with the IAS using REST. The registry is complemented by a custom VSCC that runs on all peers and that verifies the attestation report returned by the IAS. The enclave transaction validator is implemented in the form of a custom VSCC. It verifies the SGX-specific signatures on the response for the transactions produced by the chaincode enclave, and obtains the corresponding public key from the enclave registry. Signatures use 256-bit ECDSA, as provided by Go v1.10 and the SGX SDK. The ledger state as well as the transaction proposals and proposal responses are encrypted by default with 128-bit AES-GCM. The AES key for encrypting the proposal is established using a Diffie-Hellman key derivation scheme available within SGX.

7.2 Auction prototype

In order to demonstrate and evaluate the overhead of the approach, we have implemented the blockchain auction mentioned earlier.

The auction chaincode runs in an enclave. A client, the *auctioneer*, creates a new auction by invoking *create(params)* at a peer, and receives in return a unique auction identifier *auction*, which is used for any subsequent interaction with this auction. The invocation specifies a name for the auction, a description of the asset, and more. When creating a new auction, the chaincode accesses the KVS (using *getState*) and ensures that no auction with the same name already exists. Then it stores *params* using *putState* and initializes a placeholder that will store the bids. After the auction has been created, it becomes *active* and remains so until the auctioneer invokes *close(auction)*.

While the auction is active, clients acting as *bidders* may submit encrypted bids to the auction by invoking *bid(val, auction)*, where *val* denotes the value the bidder wants to offer for the asset. Each bid is stored on the blockchain as a tuple of the form (*key*, *val*) = (*auction.client*, *val*).

The auctioneer may close the auction at an arbitrary time by invoking *close(auction)*. This transaction acts as the barrier described in Section 4 and writes the updated auction status to the blockchain using *putState*. Once the auctioneer sees from its ledger that the auction is closed, it invokes *evaluate(auction)*. When the chaincode enclave receives this transaction, it determines the bidder with the highest bid and issues the transfer of the asset in exchange of the value of the bid.

7.3 Experimental setup

We deploy a Fabric network with a *solo* ordering service (one trusted node) using a single channel and three peers. Each peer and the ordering service run on a separate Supermicro 5019-MR server with a 3.4GHz four-core E3-1230 V5 Intel CPU that provides SGX support. All machines are equipped with 32 GB of memory, 1 Gbps network connection, and a SATA SSD drive; they run Ubuntu Linux 16.04 LTS Server with the generic 4.13.0-32 Linux kernel. For reporting transaction throughput we use an increasing number of clients build with the Fabric Client SDK for Go (https://github.com/hyperledger/fabric-sdk-go) invoking concurrent transaction over a period of at least 30 seconds and report the average.

As a *baseline* for our experiments we run the same auction chaincode written in Go in an unprotected environment, executed by an unmodified Fabric peer. For comparison this chaincode also uses 128-bit AES-GCM encryption to seal bids and to encrypt the auction state. However, since the peer knows the key, this does not hide the auction data from the peer.

7.4 Measurements

TCB Size. The trusted computing base (TCB) of our prototype includes the chaincode enclave and the ledger enclave; all other components of the peer are considered to be untrusted. Taken together, the system consists of approximately 5,000 lines of trusted C/C++ code and 4,000 lines of untrusted C/C++ and Go code. The trusted ledger enclave makes up the majority of the code base, with roughly 3,800 lines, whereas the chaincode enclave only comprises about 1,200 lines. The auction chaincode itself is 200 lines of C/C++. In contrast to a solution where the entire Fabric peer ($\geq 100,000$ lines) or the entire Linux kernel ($\geq 25M$ lines) is executed in the trusted environment, our approach clearly fulfills the goal of minimizing the TCB. This facilitates its security analysis through code reviews or automated verification.

Transaction size. In a preliminary experiment we evaluated the transaction sizes for the auction transactions. We observed an average transaction size of 3kB for *bid()* and 3.5kB for *evaluate()* (with a readset containing 10 bids). The transactions contain a constant overhead of about 100B for a 256-bit ECDSA signature that is formatted with JSON and produced by the chaincode enclave as described in Section 5. The overhead introduced by our approach in relation to the given transaction size in Fabric remains relatively small. As also reported in [3], transaction in Fabric are large because they contain PEM-encoded certificates.

Transaction endorsement. Next we study the endorsement throughput and latency of our approach with an increasing the number of clients. In this workload we use up to 128 clients, which concurrently invoke transactions at a single endorsement peer. Each client invokes *noop* and *submit* transactions in a closed loop, respectively. The auction chaincode returns immediately for a *noop*. The *submit* transaction, on the other hand, receives a bid, encrypts it, and stores it on the blockchain. We compare the auction chaincode executed in the SGX enclave with the native, unprotected execution.

As Fig. 5 shows, the throughput and the latency of the *noop* transaction behave almost identically. The SGX-based approach follows the baseline and scales almost linearly until reaching saturation at

Figure 5: Endorsement throughput with different numbers of clients.

	mean	σ_t	t_{min}	t_{max}
Decrypt tx	0.20	0.04	0.17	0.38
getState	0.37	0.23	0.12	4.54
cc2cc	2.59	1.42	1.08	11.44
Ledger enclave	0.68	0.16	0.52	1.42
Decrypt & verify state	0.06	0.02	0.04	0.38
Sign response	0.226	0.045	0.179	0.402

Table 1: Endorsement latency breakdown for *submit* transaction with 4 clients showing the average response time in milliseconds.

16 clients. We observe that for the *noop* transaction our approach reaches 0.93x–0.99x the throughput of the native execution. On the other hand, when executing *submit* transactions execution in SGX reaches 0.55x–95x the throughput of the native execution. In particular, before saturation our approach shows an increased latency of about 6ms compared to the native execution. The reason is that the auction chaincode invokes *getState* to retrieve data from the KVS and additionally fetches the corresponding integrity metadata from the ledger enclave using cc2cc invocations. We profiled the response latency and present the breakdown in Table 1. The table shows that retrieving data from the KVS takes 0.37ms, whereas cc2cc invocation takes 2.59ms to return. The actual invocation of the ledger enclave takes less than a millisecond, thus the majority of the response time is spent for the communication between the two chaincodes. In this experiment we see that cc2cc invocations are relatively expensive. Our choice of implementing the ledger enclave as a chaincode allowed for a simple integration and coupling with the chaincode enclave, but it comes with noticeable overhead. This overhead can be reduced by moving the ledger enclave into the peer itself, so that it directly provides Fabric's chaincode API.

Auction evaluation. We also investigate the performance of range queries used by the auction chaincode to read all submitted bids for determining the winner. We measure the response latency for the *evaluate* transaction for different numbers of submitted bids, shown in Fig. 6. As expected, we observe that latency increases with larger number of submitted bids. The relative overhead of executing in an SGX enclave remains constant for a small number of submitted bids, at about 20%, and for larger numbers it decreases to 10%. This experiment shows that using range queries reduces the number of cc2cc invocations, which improves the performance.


Figure 6: Auction evaluation latency with different number of submitted bids.

End-to-End performance. In the last experiment we study overall end-to-end response latency for the auction. The workload is the same as in the transaction endorsement experiment but here we measure the throughput and latency for the entire transaction flow including endorsement, ordering, and validation. We use the default block size configuration with 10 transactions per block. We observe that execution in SGX reaches 0.80x–0.95x of the throughput achieved by the native execution, as shown in Fig. 7.


Figure 7: End-to-end throughput with different numbers of clients.

8 RELATED WORK

Trusted execution technology is envisaged to play an important role in the context of blockchains, especially for enterprise applications and in consortium blockchains. The two most prominent uses of TEEs are to execute smart contracts for keeping data private and to represent off-chain data securely. We review these two and further related applications of TEEs in this section.

Smart-contract execution with Intel SGX. Several approaches have recently suggested to execute blockchain applications and smart contracts within Intel SGX.

The most prominent among these and most closely related to our work is the *Coco framework*, announced by Microsoft in a white paper [32]. It provides a set of building blocks based on Intel SGX that can be used to secure blockchain systems. Coco integrates consensus algorithms, distributed ledger state, and a runtime environment for executing smart-contract transactions in SGX enclaves. It appears to be derived from Ethereum and mentions the Ethereum Virtual Machine (EVM) for its core data structures and protocols. The components of Coco are described as potentially separate enclaves, but conceptually the entire blockchain node (corresponding to a peer in Fabric) resides in SGX. Coco exploits concepts from a proposed related blockchain system called *VOLT* [40]. Since only a white paper for Coco is available, however, it is difficult to clearly assess the framework. It is clear that Coco suffers from a large TCB with its problems as discussed in Section 4.3, including potential malicious interactions across smart contracts. Our approach isolates each application within its own enclave with the minimal amount of code necessary and follows the philosophy of minimizing the TCB.

The *R3 Corda* distributed ledger platform has also announced a privacy feature using SGX in a white paper [17]. In Corda, some aspects of the transaction validation are envisaged to take place inside an SGX enclave, potentially running at untrusted nodes. By executing the transaction validation in an enclave, the identities involved in a transaction can be encrypted on the blockchain and are only revealed inside an enclave during validation. Corda strives to port a Java runtime environment (JRE) in order to execute native Corda smart contracts within an enclave. Compared to our approach, this introduces some runtime overhead for the JRE and an increased attack surface with the larger TCB, but it benefits from portability of the applications written in the same language.

The *IBM Blockchain Platform* [20] offers an enterprise blockchain-as-a-service solution allowing for deployment of a Fabric network using Secure Service Container (SSC) technology [19] on *IBM Z* systems. The platform runs the whole peer and its Linux operating system within a secure container, which is shielded from access of the host and host administrator, similar to SGX. This means that running Fabric within a Secure Service Container requires specialized mainframe hardware and has a large TCB. In contrast, our approach for running Fabric with SGX works with commodity systems and minimizes the TCB for each smart-contract application.

Several academic papers have also suggested to run smart contracts inside SGX for confidentiality, e.g., in the "Ring of Gyges" [23] or in Hawk [26].

Blockchain oracles and off-chain data. Other works [33, 47] realize trusted "oracles" for blockchain smart contracts using SGX. Oracles are data feeds external to the blockchain that inform a smart contract about "facts" in the environment. They extend the scope of inputs to which an application can respond and serve as trusted sources and triggers for actions on the blockchain. Leveraging trusted execution technology enhances the trustworthiness of an oracle and allows to verify the correctness of the data source. This work is orthogonal to our approach, which could also benefit from oracles that exploit trusted hardware.

Teechain [28] is a system to perform off-chain payments on top of Bitcoin. It leverages SGX to establish stateful payment channels among mistrusting parties. Such off-chain channels resolve exchanges bilaterally without incurring a blockchain transaction for every exchange, in the normal case when both parties are honest. Payment channels expected to boost the overall throughput of a blockchain-based payment system. Through the use of SGX, Teechain relaxes the synchrony assumptions in existing payment channels and gains efficiency.

Consensus protocols. Another line of work leverages trusted execution to enhance the resilience and performance of consensus protocols. Based on traditional BFT protocols, systems such as *TrInc* [27], *CheapBFT* [24], or Hybster [8] have shown how to enhance state-machine replication with trusted specialized hardware devices, FPGAs, and SGX enclaves, respectively.

Some blockchain-specific peer-to-peer consensus protocols have been introduced proposed that scale to a large number of nodes based on trusted hardware. *REM* [48], for example, introduces *Proofs-of-Useful-Work* that are run within SGX in order to reach consensus on a public blockchain.

Furthermore, the consensus model of the Hyperledger Sawtooth platform (https://github.com/hyperledger/sawtooth-core), originally contributed by Intel, includes the *Proof-of-Elapsed Time (PoET)* consensus protocol based on SGX. It replaces the proof-of-work function for leader election in Bitcoin's Nakamoto consensus with a mandatory, random waiting time imposed by an enclave running in SGX; otherwise the protocol is similar to the Nakamoto consensus.

State continuity and TEEs. Kaptchuk et al. [25] address state continuity for memoryless secure processors that have access to a distributed ledger. They construct a generic protocol for detecting rollback attacks, assuming the processor is always given access to latest ledger state. This is an interesting conceptual approach to detect rollbacks, assuming an idealized trusted ledger that cannot be rolled back.

9 CONCLUSION

This work has explored some pitfalls that arise from the combination of trusted execution with blockchains. In particular, smart-contract execution with Intel SGX promises protection for blockchain applications with strong privacy demands. However, since enclaves are susceptible to rollback attacks, a malicious blockchain peer may break confidentiality by resetting an enclave to manipulate the operation invocation order. We have presented a solution that selectively utilizes SGX, minimizes the TCB, and handles rollbacks for the Hyperledger Fabric platform. An evaluation has shown that the overhead of our approach is within 10%–20% for a sealed-bid auction application.

REFERENCES

[1] Joy Algesheimer, Christian Cachin, Jan Camenisch, and Günter Karjoth. 2001. Cryptographic Security for Mobile Code. In *Proc. IEEE Symposium on Security and Privacy*.

[2] Ittai Anati, Shay Gueron, Simon Johnson, and Vincent Scarlata. 2013. Innovative Technology for CPU Based Attestation and Sealing. In *Int. Workshop on Hardware and Architectural Support for Security and Privacy (HASP)*. ACM.

[3] Elli Androulaki, Artem Barger, Vita Bortnikov, Christian Cachin, Konstantinos Christidis, Angelo De Caro, David Enyeart, Christopher Ferris, Gennady Laventman, Yacov Manevich, Srinivasan Muralidharan, Chet Murthy, Binh Nguyen, Manish Sethi, Gari Singh, Keith Smith, Alessandro Sorniotti, Chrysoula Stathakopoulou, Marko Vukolić, Sharon Weed, and Jason Yellick. 2018. Hyperledger Fabric: A Distributed Operating System for Permissioned Blockchains. In *Proc. 13th European Conference on Computer Systems (EuroSys)*.

[4] Maria Apostolaki, Aviv Zohar, and Laurent Vanbever. 2017. Hijacking Bitcoin: Routing Attacks on Cryptocurrencies. In *Proc. 38th IEEE Symposium on Security & Privacy*.

[5] Sergei Arnautov, Bohdan Trach, Franz Gregor, Thomas Knauth, Andre Martin, Christian Priebe, Joshua Lind, Divya Muthukumaran, Dan O'Keeffe, Mark L. Stillwell, David Goltzsche, Dave Eyers, Rüdiger Kapitza, Peter Pietzuch, and Christof Fetzer. 2016. SCONE: Secure Linux Containers with Intel SGX. In *USENIX Proc. on Operating Systems Design and Implementation (OSDI)*.

[6] Artem Barger, Yacov Manevich, Benjamin Mandler, Vita Bortnikov, Gennady Laventman, and Gregory Chockler. 2017. Scalable Communication Middleware for Permissioned Distributed Ledgers. In *Proc. 10th ACM Int. Systems and Storage Conference (SYSTOR)*. ACM.

[7] Andrew Baumann, Marcus Peinado, and Galen Hunt. 2014. Shielding Applications from an Untrusted Cloud with Haven. In *USENIX Proc. on Operating Systems Design and Implementation (OSDI)*.

[8] Johannes Behl, Tobias Distler, and Rüdiger Kapitza. 2017. Hybrids on Steroids: SGX-based High Performance BFT. In *Proc. 12th European Conference on Computer Systems (EuroSys)*.

[9] Marcus Brandenburger, Christian Cachin, Matthias Lorenz, and Rüdiger Kapitza. 2017. Rollback and Forking Detection for Trusted Execution Environments using Lightweight Collective Memory. In *Proc. 47th Int. Conference on Dependable Systems and Networks (DSN)*.

[10] Ferdinand Brasser, Urs Müller, Alexandra Dmitrienko, Kari Kostiainen, Srdjan Capkun, and Ahmad-Reza Sadeghi. 2017. Software Grand Exposure: SGX Cache Attacks Are Practical. In *Proc. USENIX Conference on Offensive Technologies (WOOT '17)*.

[11] Stefan Brenner, Colin Wulf, Matthias Lorenz, Nico Weichbrodt, David Goltzsche, Christof Fetzer, Peter Pietzuch, and Rüdiger Kapitza. 2016. SecureKeeper: Confidential ZooKeeper using Intel SGX. In *Int. Middleware Conference*. ACM.

[12] Ernie Brickell and Jiangtao Li. 2011. Enhanced privacy ID from bilinear pairing for hardware authentication and attestation. *Int. Journal of Information Privacy, Security and Integrity 2* (2011).

[13] Christian Cachin and Marko Vukolic. 2017. Blockchain Consensus Protocols in the Wild. In *Proc. 31st Intl. Symposium on Distributed Computing (DISC 2017)*.

[14] Bernardo David, Peter Gazi, Aggelos Kiayias, and Alexander Russell. 2018. Ouroboros Praos: An Adaptively-Secure, Semi-Synchronous Proof-of-Stake Blockchain. In *Advances in Cryptology: Eurocrypt 2018*. Springer.

[15] Juan A. Garay, Aggelos Kiayias, and Nikos Leonardos. 2015. The Bitcoin Backbone Protocol: Analysis and Applications. In *Advances in Cryptology: Eurocrypt 2015*. Springer.

[16] David Goltzsche, Colin Wulf, Divya Muthukumaran, Konrad Rieck, Peter Pietzuch, and Rüdiger Kapitza. 2017. TrustJS: Trusted Client-side Execution of JavaScript. In *Proc. 10th European Workshop on System Security (EuroSec)*.

[17] Mike Hearn. 2016. Corda: A distributed ledger. (2016). Whitepaper, https://docs.corda.net/_static/corda-technical-whitepaper.pdf.

[18] Matthew Hoekstra, Reshma Lal, Pradeep Pappachan, Vinay Phegade, and Juan Del Cuvillo. 2013. Using Innovative Instructions to Create Trustworthy Software Solutions. In *Int. Workshop on Hardware and Architectural Support for Security and Privacy (HASP)*. ACM.

[19] IBM. 2017. Secure Service Container User's Guide SC28-6971-01. (2017). https://www-01.ibm.com/support/docview.wss?uid=isg2bb79df265313634d85258088005188e3&aid=1.

[20] IBM. 2018. IBM Blockchain Platform. (2018). https://console.bluemix.net/docs/services/blockchain/index.html.

[21] Intel. 2018. Intel Software Guard Extensions (Intel SGX) Developer Guide. (2018). https://01.org/intel-software-guard-extensions/documentation/intel-sgx-developer-guide.

[22] Intel. 2018. Intel Software Guard Extensions (Intel SGX) SDK for Linux* OS. (2018). https://01.org/intel-software-guard-extensions/documentation/intel-sgx-sdk-developer-reference.

[23] Ari Juels, Ahmed E. Kosba, and Elaine Shi. 2016. The Ring of Gyges: Investigating the Future of Criminal Smart Contracts. In *Proc. ACM SIGSAC Conference on Computer and Communications Security (CCS)*.

[24] Rüdiger Kapitza, Johannes Behl, Christian Cachin, Tobias Distler, Simon Kuhnle, Seyed Vahid Mohammadi, Wolfgang Schröder-Preikschat, and Klaus Stengel. 2012. CheapBFT: Resource-efficient Byzantine Fault Tolerance. In *Proc. 7th European Conference on Computer Systems (EuroSys)*.

[25] Gabriel Kaptchuk, Ian Miers, and Matthew Green. 2017. Giving State to the Stateless: Augmenting Trustworthy Computation with Ledgers. (2017). Cryptology ePrint Archive, Report 2017/201.

[26] Ahmed E. Kosba, Andrew Miller, Elaine Shi, Zikai Wen, and Charalampos Papamanthou. 2016. Hawk: The Blockchain Model of Cryptography and Privacy-Preserving Smart Contracts. In *Proc. 37th IEEE Symposium on Security & Privacy*.

[27] Dave Levin, John R. Douceur, Jacob R. Lorch, and Thomas Moscibroda. 2009. TrInc: Small Trusted Hardware for Large Distributed Systems.. In *Proc. Networked Systems Design and Implementation (NSDI)*.

[28] Joshua Lind, Ittay Eyal, Florian Kelbert, Oded Naor, Peter R. Pietzuch, and Emin Gün Sirer. 2017. Teechain: Scalable Blockchain Payments using Trusted Execution Environments. *CoRR* (2017). arXiv:1707.05454 http://arxiv.org/abs/1707.05454

[29] Joshua Lind, Christian Priebe, Divya Muthukumaran, Dan O'Keeffe, Pierre-Louis Aublin, Florian Kelbert, Tobias Reiher, David Goltzsche, David Eyers, Rüdiger Kapitza, Christof Fetzer, and Peter Pietzuch. 2017. Glamdring: Automatic Application Partitioning for Intel SGX. In *USENIX Annual Technical Conference (USENIX ATC)*.

[30] Sinisa Matetic, Mansoor Ahmed, Kari Kostiainen, Aritra Dhar, David Sommer, Arthur Gervais, Ari Juels, and Srdjan Capkun. 2017. ROTE: Rollback Protection for Trusted Execution. In *Proc. USENIX Security Symposium*.

[31] Frank McKeen, Ilya Alexandrovich, Alex Berenzon, Carlos V. Rozas, Hisham Shafi, Vedvyas Shanbhogue, and Uday R. Savagaonkar. 2013. Innovative Instructions and Software Model for Isolated Execution. In *Int. Workshop on Hardware and Architectural Support for Security and Privacy (HASP)*. ACM.

[32] Microsoft. 2017. The Coco Framework. (2017). Whitepaper, https://github.com/Azure/coco-framework.

[33] Microsoft. 2017. Introducing Project "Bletchley". (2017). Whitepaper, https://github.com/Azure/azure-blockchain-projects/blob/master/bletchley/bletchley-whitepaper.md.

[34] Paul Milgrom. 2004. *Putting Auction Theory to Work*. Cambridge University Press.

[35] Satoshi Nakamoto. 2009. Bitcoin: A Peer-to-Peer Electronic Cash System. Whitepaper. (2009). http://bitcoin.org/bitcoin.pdf.

[36] Meni Orenbach, Pavel Lifshits, Marina Minkin, and Mark Silberstein. 2017. Eleos: ExitLess OS Services for SGX Enclaves. In *Proc. 12th European Conference on Computer Systems (EuroSys)*.

[37] Bryan Parno, Jacob R Lorch, John R Douceur, James Mickens, and Jonathan M McCune. 2011. Memoir: Practical state continuity for protected modules. In *Proc. IEEE Symposium on Security & Privacy*.

[38] Fred B. Schneider. 1990. Implementing Fault-Tolerant Services using the State Machine Approach: A Tutorial. *Comput. Surveys* 22, 4 (Dec. 1990), 299–319.

[39] Michael Schwarz, Samuel Weiser, Daniel Gruss, Clémentine Maurice, and Stefan Mangard. 2017. Malware guard extension: Using SGX to conceal cache attacks. In *Int. Conference on Detection of Intrusions and Malware, and Vulnerability Assessment*. Springer.

[40] Srinath Setty, Soumya Basu, Lidong Zhou, Michael Lowell Roberts, and Ramarathnam Venkatesan. 2017. *Enabling secure and resource-efficient blockchain networks with VOLT*. Technical Report. https://www.microsoft.com/en-us/research/publication/enabling-secure-resource-efficient-blockchain-networks-volt/

[41] Ming-Wei Shih, Sangho Lee, Taesoo Kim, and Marcus Peinado. 2017. T-SGX: Eradicating controlled-channel attacks against enclave programs. In *Proc. 2017 Annual Network and Distributed System Security Symposium (NDSS)*.

[42] Raoul Strackx, Bart Jacobs, and Frank Piessens. 2014. ICE: A Passive, High-Speed, State-Continuity Scheme. In *Annual Computer Security Applications Conference*. ACM.

[43] Raoul Strackx and Frank Piessens. 2016. Ariadne: A minimal approach to state continuity. In *USENIX Security Symposium*.

[44] Chia-che Tsai, Donald E. Porter, and Mona Vij. 2017. Graphene-SGX: A Practical Library OS for Unmodified Applications on SGX. In *USENIX Annual Technical Conference (USENIX ATC)*.

[45] Gavin Wood. 2016. Ethereum: A Secure Decentralised Generalised Transaction Ledger. http://gavwood.com/paper.pdf. (2016).

[46] Yuanzhong Xu, Weidong Cui, and Marcus Peinado. 2015. Controlled-channel attacks: Deterministic side channels for untrusted operating systems. In *Proc. IEEE Symposium on Security and Privacy*.

[47] Fan Zhang, Ethan Cecchetti, Kyle Croman, Ari Juels, and Elaine Shi. 2016. Town Crier: An Authenticated Data Feed for Smart Contracts. In *Proc. ACM SIGSAC Conference on Computer and Communications Security (CCS)*.

[48] Fan Zhang, Ittay Eyal, Robert Escriva, Ari Juels, and Robbert van Renesse. 2017. REM: Resource-Efficient Mining for Blockchains. In *Proc. 26th USENIX Security Symposium*.